DL2023

Instructor: Dr. Mayank Vatsa

HUMAN FACE EMOTION RECOGNITION

DL LAB PROJECT REPORT

By-

Riyanshu Jain (B20AI060)

Ruthvik K (B20AI037)

Shubh Soni (B20BB039)

Abstract

Using the FER2013 dataset, we developed a method for recognising human facial expressions. The dataset includes images labelled with seven different emotion classes. To identify the best deep learning model for this task, we trained and compared the performance of several deep learning models, including custom model architecture, ResNet18, MobileNetV2, SqueezeNet1_0, and ShuffleNetV2_x1_0.

Introduction

Because of wide range of applications, such as surveillance, human-computer interaction, and medical diagnostics, human face recognition has been an active area of research in computer vision and machine learning. The accurate identification of facial expressions, which convey information about emotions, is one challenge in face recognition. The FER2013 dataset, which contains 48x48 pixel grayscale images of faces labelled with seven emotion categories, is a widely used dataset for facial expression recognition tasks.

Methods

We used PyTorch in this project to implement, train, and test various deep learning models. The dataset was divided into two parts: a training set and a test set, with data augmentation techniques like random cropping, rotation, and horizontal flipping applied to the training set to improve the model's generalisation capabilities. We tried out the following models:

- Custom Model: A custom CNN architecture was implemented, consisting of four convolutional layers followed by max-pooling and dropout layers, and two fully connected layers.
- ResNet18: A pre-trained ResNet18 model was fine-tuned for the facial expression recognition task.
- MobileNetV2: A pre-trained MobileNetV2 model was fine-tuned for the facial expression recognition task.
- SqueezeNet1_0: A pre-trained SqueezeNet1_0 model was fine-tuned for the facial expression recognition task.
- ShuffleNetV2_x1_0: A pre-trained ShuffleNetV2_x1_0 model was fine-tuned for the facial expression recognition task.

For each model, the optimizer used was Adam, and the loss function was CrossEntropyLoss. The models were trained for 50 epochs.

Results

After training and testing the models, the following accuracies were achieved:

1. Custom Model: Accuracy: 62.76%
2. ResNet18: Accuracy: 67.54%
3. MobileNetV2: Accuracy: 66.95%
4. SqueezeNet1_0: Accuracy: 59.07%
5. ShuffleNetV2_x1_0: Accuracy: 67.69%

Confusion matrices were plotted to visualize the performance of each model across the different emotion classes.

Custom Model

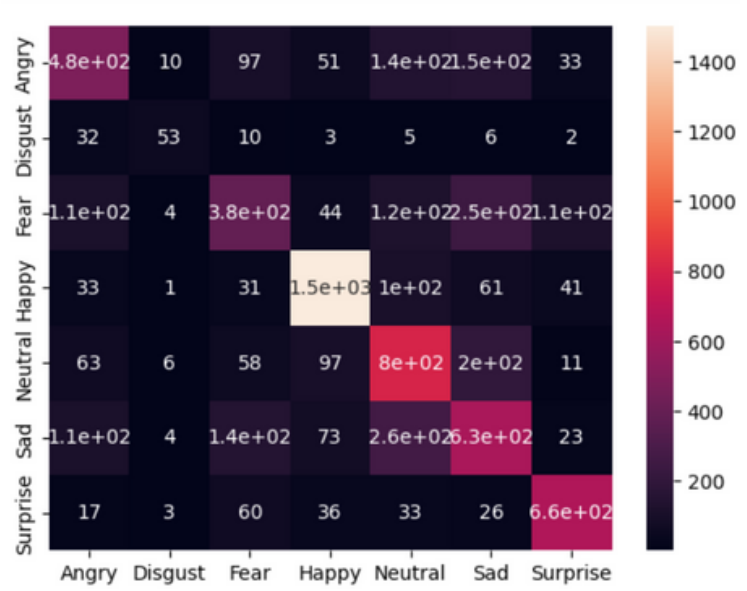

ResNet Model

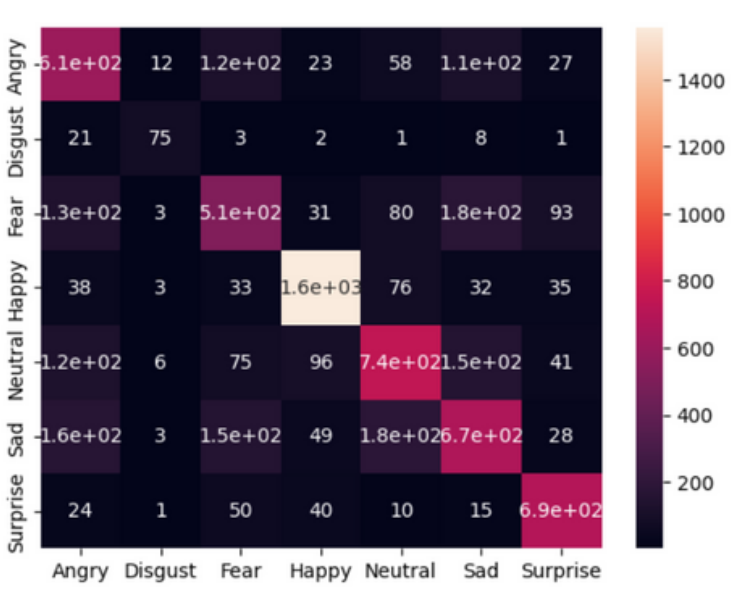

Mobilenet Model

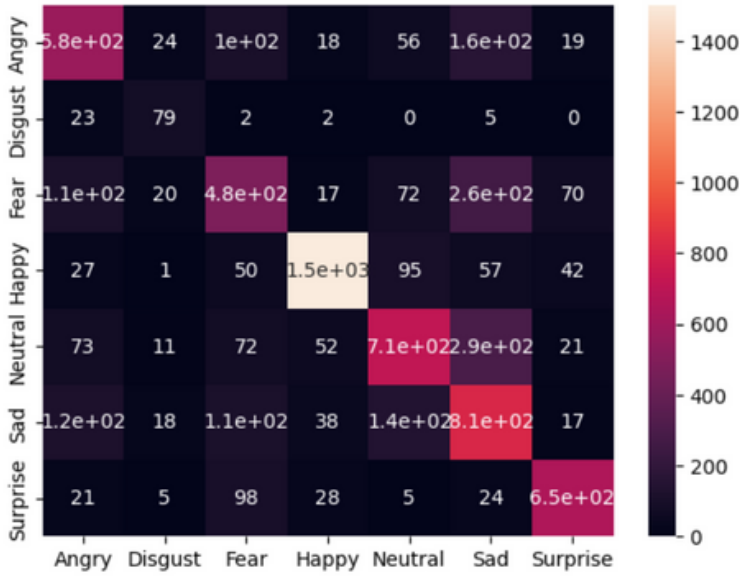

Squeezenet Model

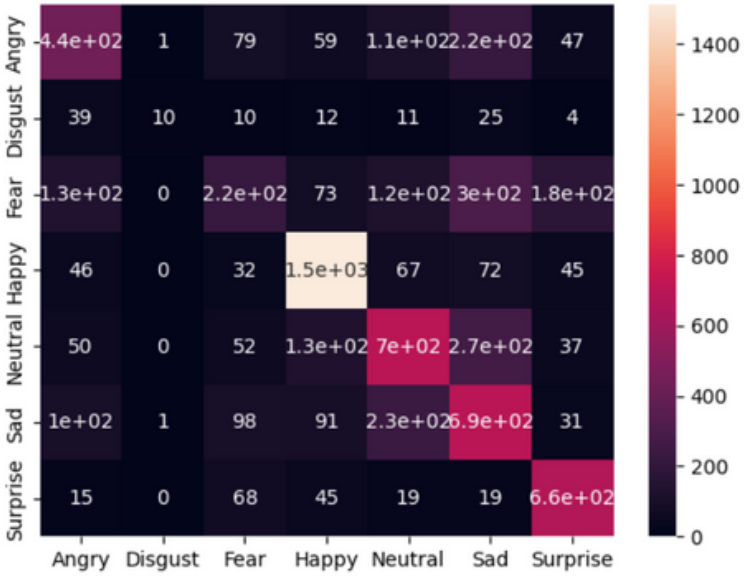

Shufflenet Model

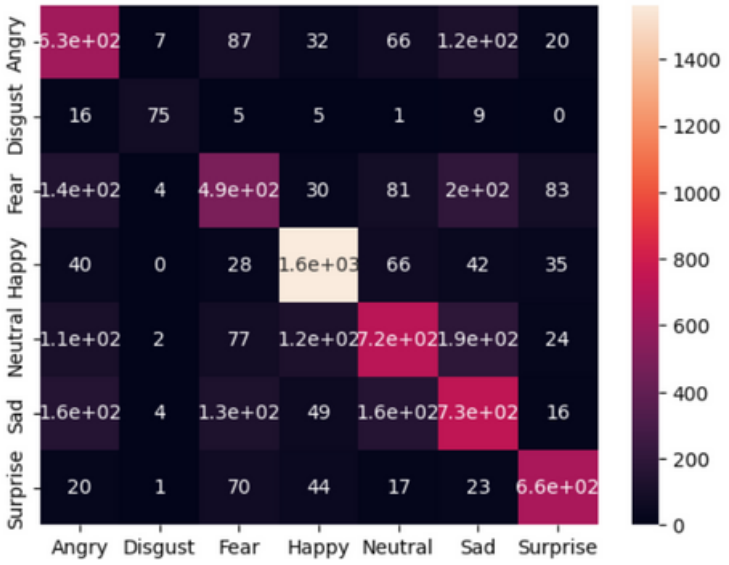

TensorBoard Visulisation Of Each Model

Resnet Model
Average Step Time
496,908

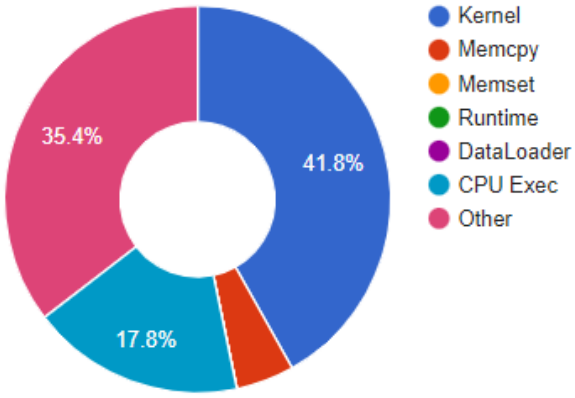

Mobilenet Model
Average Step Time
496,908

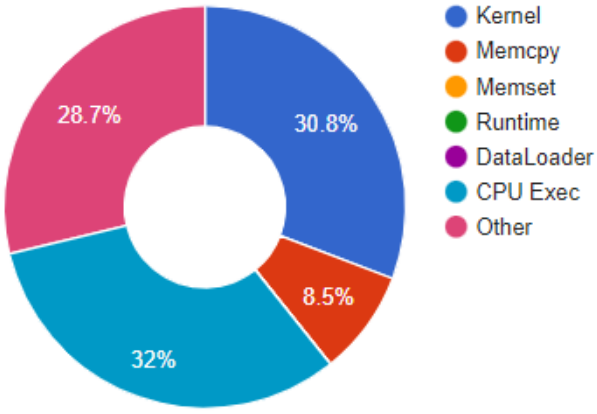

Squeezenet Model
Average Step Time
469,653

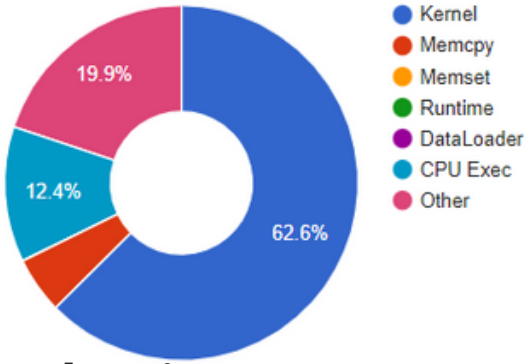

Shufflenet Model
Average Step Time
988,294

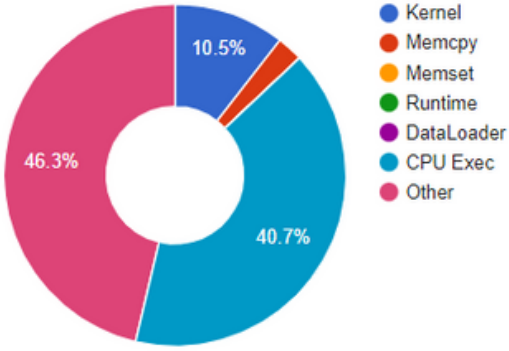

Conclusion

Using the FER2013 dataset, we investigated various deep learning models for human face recognition in this project. The results show that the ResNet18 model is effective at recognising facial expressions as it have higher accuracy and lower average step time. Future research should concentrate on improving model architectures and using larger datasets to improve the performance of facial expression recognition models.

Deploying our Model

- We have first deployed our model on local host using Flask later using StreamLit we have deployed our model on Huggingface.

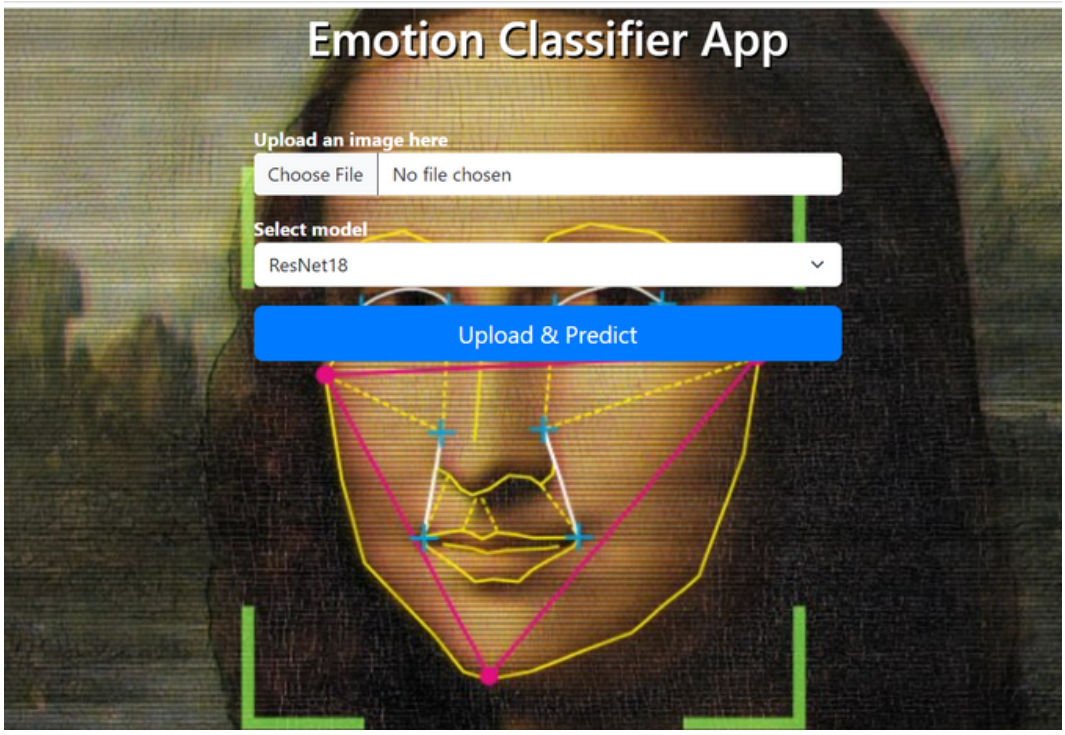

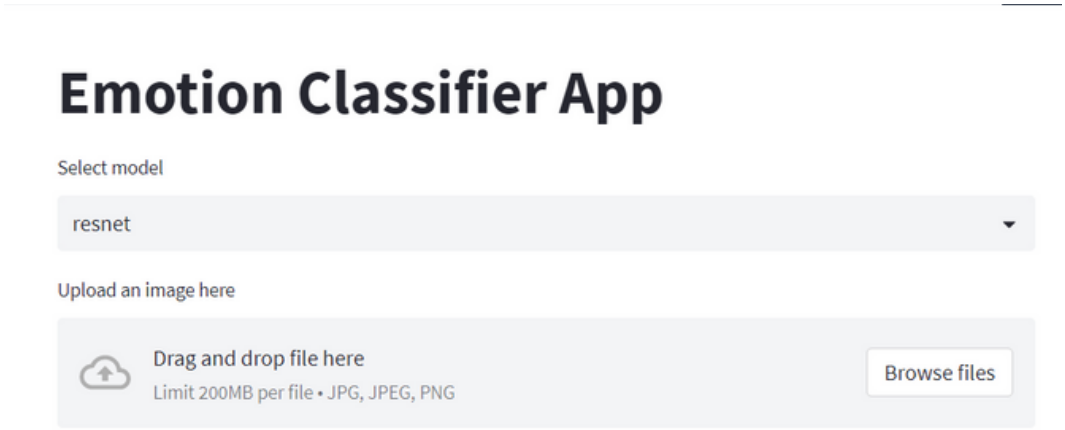

- Link where the model is deployed using HuggingFace

https://huggingface.co/spaces/ruthvik2/DLOps_Project_1